

# CATMA: Conformance Analysis Tool For Microservice Applications

Clinton Cao  
Delft University of Technology  
The Netherlands

Simon Schneider  
Hamburg University of Technology  
Germany

Nicolás E. Díaz Ferreyra  
Hamburg University of Technology  
Germany

Sicco Verwer  
Delft University of Technology  
The Netherlands

Annibale Panichella  
Delft University of Technology  
The Netherlands

Riccardo Scandariato  
Hamburg University of Technology  
Germany

## ABSTRACT

The microservice architecture allows developers to divide the core functionality of their software system into multiple smaller services. However, this architectural style also makes it harder for developers to debug and assess whether the system's deployment conforms to its implementation. We present CATMA, an automated tool that detects non-conformances between the system's deployment and implementation. It automatically visualizes and generates potential interpretations for the detected discrepancies. Our evaluation of CATMA shows promising results in terms of performance and providing useful insights. CATMA is available at <https://github.com/tudelft-cda-lab/catma>, and a demonstration video is available at <https://youtu.be/LxRDPjRa0l8>.

## CCS CONCEPTS

• **Software and its engineering** → **Software testing and debugging**; **Automated static analysis**; **Dynamic analysis**.

## KEYWORDS

microservices, static analysis, dynamic analysis, software testing, empirical software engineering

### ACM Reference Format:

Clinton Cao, Simon Schneider, Nicolás E. Díaz Ferreyra, Sicco Verwer, Annibale Panichella, and Riccardo Scandariato. 2018. CATMA: Conformance Analysis Tool For Microservice Applications. In *Proceedings of International Conference on Software Engineering (ICSE'24)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/XXXXXXX.XXXXXXX>

## 1 INTRODUCTION

Software systems following the microservice architectural paradigm have their core functionality split into multiple smaller components. These microservices (or just *services*) of a microservice application (MSA) communicate via lightweight communication protocols such as REST APIs or message brokers. The services of an MSA can be developed, maintained, and deployed independently,

paving the way for an increasing trend in the adoption of this architectural style. Despite these benefits, MSAs pose a challenge in gaining a comprehensive overview due to their inherently decoupled and distributed nature. Consequently, debugging faults is a time-consuming process because the localization of the root cause is challenging. According to studies, developers usually take several days to debug and find the cause of a fault [8, 16]. Many approaches for the automatic extraction of an architectural representation of MSA have been proposed [1, 5, 9, 12], thus addressing the challenge of gaining an overview. Some approaches combine static and dynamic analysis to build the models. Also, multiple fault localization techniques for MSAs have been proposed [6, 17], which use dynamic analysis to identify faults and pinpoint the root cause in code. However, no work compares the results from static and dynamic analysis rather than merging them. Also, none of the fault localization approaches offer explainability in the form of possible interpretations for the faults.

In this paper, we present CATMA, a pioneering tool designed to analyze and compare statically and dynamically obtained architectural models. CATMA autonomously identifies potential non-conformances between these models, generating easily accessible visualizations for users and providing concise interpretations. These interpretations reduce the number of lines in source code that users need to scrutinize when investigating a non-conformance. We tested CATMA on four open-source MSA and conducted a preliminary usability study with two participants. The results indicate that the tool effectively supports developers during the localization and debugging of non-conformances, demonstrating its usefulness and potential in the debugging landscape for microservices.

## 2 RUNNING EXAMPLE

The software engineering team of ZYX Inc. is working on their new web application for selling tech products. They decide to develop their application following the microservice architectural style as this allows them to split up into smaller groups and work independently on the core functionalities of their application. Each member follows the best practices of software engineering; using static analysis to detect faults and testing each functionality before its deployment. After finishing the development, they deploy the application to test it out. To their surprise, they notice that the monitoring service does not receive any metrics data. They are unsure of the cause of this discrepancy since a static analysis tool correctly detects the line of code that implements the sending of metrics data and does not raise any warnings. They spend several


Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

ICSE'24, April 14–20, 2024, Lisbon, Portugal

© 2018 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/XXXXXXX.XXXXXXX>



**Figure 1: CATMA’s workflow.** Input models are parsed (①) and non-conformances are detected between them (②) and visualized (④). For each non-conformance, possible interpretations are generated (③) and visualized (⑤).

days analyzing different log files, but have no luck in finding the underlying cause. They scratch their heads and start wondering whether there is a tool that provides the following functionalities:

- Detecting discrepancies between the implementation and deployment of an arbitrary MSA,
- Providing a high-level overview of detected discrepancies, and
- Providing textual descriptions of the potential causes for detected discrepancies.

### 3 CATMA

**Workflow.** Figure 1 depicts CATMA’s workflow. First, the *Model Processor* ① parses the input models (static and dynamic) to extract architectural components. The obtained data is passed on to the *Non-conformance Detector* ②, which checks whether there are any discrepancies between the two input models. If a non-conformance is detected, it is forwarded to both the *Interpretation Generator* ③ and the *Non-conformance Visualizer* ④. The latter (④) collects all detected non-conformances and generates a model-based visualization, showing the non-conformances in the system’s architecture. The former (③) generates a set of possible interpretations for each detected non-conformance, which describe potential causes. These interpretations are forwarded to the *Interpretation Visualizer* ⑤, which generates HTML pages that visualize the interpretations. CATMA is designed to be modular, meaning that each component presented in the workflow can be replaced or expanded to fit its user’s needs.

**Detecting Non-conformances.** As static models, CATMA accepts dataflow diagrams (DFDs) extracted by the automatic approach presented by Schneider and Scandariato [11]. As dynamic models, state machines inferred from HTTP events logs are expected. They are created using a similar model-inference approach as presented by Cao et al. [3]. The *Model Processor* extracts services and their interrelations from both input models. In the DFD, nodes and edges depict the services and their corresponding relations, respectively. We can therefore directly extract the nodes and edges. In a state machine, services and their corresponding relations are represented differently; each transition in a state machine indicates which services in the system have communicated with each other. Thus, nodes and edges are extracted from the transitions of the state machines. The *Model Processor* creates a set of nodes and edges for both input models, where edges are represented as “service  $X \rightarrow$  service  $Y$ ” and denote the communication relationship between the two services.

Non-conformances are detected by identifying differences between the sets of nodes and edges. The *Non-conformance Detector*

iterates through the sets and checks for each item whether it is present in both corresponding sets. Each item is tagged according to this comparison, i.e., indicating whether it is present in both, only the static, or only the dynamic model. The enriched sets of nodes and edges are passed to components ③ and ④.

The *Non-conformance Visualizer* creates a graphical depiction of possibly found non-conformances. It generates a PlantUML <sup>1</sup> file which presents the nodes and edges as a graph and where any found non-conformances are highlighted by a coloring scheme. Model items observed in both models are colored black, items only observed in the static model are colored green, and items only observed in the dynamic model are colored red.

**Interpreting Non-conformances.** CATMA generates a set of possible interpretations for each detected non-conformance. These interpretations are visualized in an HTML page by the *Interpretation Visualizer*. The HTML page helps users analyze the potential causes of non-conformance.

We define *static non-conformances* as nodes or edges only detected in the static model and *dynamic non-conformances* as those only detected in the dynamic model. CATMA generates a specific set of interpretations for each type of non-conformance. However, the generated HTML pages follow the same structure consisting of (1) the type and definition of the non-conformance and the services involved in the detected non-conformance, (2) the set of possible interpretations, and (3) additional details that support the understanding of the detected non-conformance. Parts 2 and 3 are presented individually in the following.

**Providing Interpretations of Non-Conformances.** A set of high-level textual interpretations is provided, which describe possible underlying causes of the detected non-conformances. The interpretations are meant to serve as possible starting points to debug found non-conformances. Currently, the generation is based solely on the type of non-conformance, i.e., whether it is static or dynamic. We formulated a text describing possible interpretations for both types of non-conformances, and the corresponding one is presented to the user. As the basis for these interpretations, we collected known causes of non-conformances from the literature [2, 7, 14, 15]. These causes range from standard programming errors made in software development to common causes for issues encountered by developers of MSAs. As an example, misconfiguration of services is a common cause of dynamic non-conformances in MSAs. When services are not properly configured, they become undiscoverable by other services, leading to missing expected runtime behaviors. CATMA uses this information as a basis for the generation of one interpretation for a dynamic non-conformance. For the collection, we disregarded non-conformances that are based on hardware, e.g., due to non-deterministic behavior because of multi-threading or similar effects. The textual descriptions of possible interpretations provided for a static non-conformance are shown in Figure 2. Our future work will predominantly focus on this part of the tool, specifically on implementing a more intelligent generation of applicable interpretations. In this regard, we will analyze indicators for each cause of non-conformances. These indicators will then be used to decide whether a cause is plausible or not for a given non-conformance. This will lead to the generation of a tailored set

<sup>1</sup>plantuml.com

## Potential Interpretations For the Non-Conformance

### Implicit Call via Third-Party Services

Communication flow has been detected between the two services but the responsible line of code is not detected in the source code. It could be that the call is implicitly triggered by a third-party service that is used in the implementation (e.g. via an annotation that is used by a framework or that code is injected during run-time).

### Unintentional Endpoint Exposure

The communication flow detected between the two services could be caused by an endpoint that is exposed unintentionally by the developer. It could be the case that code has been refactored and the endpoint was not removed during refactoring, or that the endpoint was used for testing purposes and was not removed afterwards.

### Code Located Outside of Default Source Location

The line of code that is responsible for triggering the flow of communication between the two services is not located in the default source-code folder of the project. It could be the case that the line of code was unintentionally introduced in a different folder (e.g. resources folder).

Figure 2: Example set of textual interpretations.

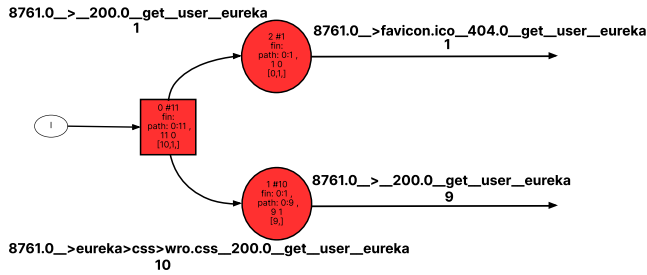


Figure 3: Part of state machine modeling unexpected behavior.

of possible interpretations for each found non-conformance. The already carried-out analysis of the related literature provides the basis for this future work.

**Additional details.** The generated HTML page also presents additional details that could aid the user with the understanding of the detected non-conformance. In the case of static non-conformances, a state machine is visualized that depicts the unexpected sequential communication behavior detected between the involved services. The most frequently occurring calls between the involved services are presented in a human-readable format right after the state machine model. This insight can be used to understand why such calls were made between the involved services. Figures 3 and 4 show an example of a state machine and the most frequent calls, respectively. In the case of a dynamic non-conformance, we instead leverage the traceability information contained in the static model to point to the code that shows the expected behavior. Specifically, the page presents (1) the line of source code responsible for triggering the missing runtime event, (2) the sequence of events that should trigger the missing runtime event, and (3) human-readable call details extracted for the previous point. Figure 5 shows an example of a part of this set of details. Furthermore, the state machines learned for each involved service are presented at the end of the HTML page.

## 4 TOOL EVALUATION

**Performance Analysis.** We evaluated CATMA's performance in terms of time to detect non-conformances in MSA. For this, we selected 4 DFDs of open-source MSAs from the dataset created by Schneider et al.[10], deployed these MSA, and created state machine models for them. Then, we ran CATMA on the obtained models and measured the time of the analysis. Table 1 presents the time

## Frequently occurring endpoint calls extracted from the dynamic model learned for the link between user and eureka:

### Endpoint: /

- Port: 8761.0
- Call status code: 200.0
- Call direction: from user to eureka
- Call frequency: 38

### Endpoint: /eureka/css/wro.css

- Port: 8761.0
- Call status code: 200.0
- Call direction: from user to eureka
- Call frequency: 10

Figure 4: Most frequent calls for unexpected behavior.

Sequences that occurred in the dynamic model that should produce run-time behaviour for link between order and turbine
<ul style="list-style-type: none"> <li>• user →(implicit) zuul → order → turbine</li> <li>• zuul → order → turbine</li> </ul>
For the occurred sequences, these are the unique sequence of endpoints (parameters) that were used in the sequence
<ul style="list-style-type: none"> <li>• Sequence: user → zuul → order → turbine <ul style="list-style-type: none"> <li>◦ Call started with "/order/line". Then followed by call with "/line".</li> <li>◦ Call started with "/order/". Then followed by call with "/".</li> <li>◦ Call started with "/order/18". Then followed by call with "/18".</li> <li>◦ Call started with "/order/11". Then followed by call with "/11".</li> </ul> </li> <li>• Sequence: zuul → order → turbine <ul style="list-style-type: none"> <li>◦ Call started with "/line".</li> <li>◦ Call started with "/18".</li> <li>◦ Call started with "/3".</li> <li>◦ Call started with "/5".</li> <li>◦ Call started with "/11".</li> <li>◦ Call started with "/form.html".</li> <li>◦ Call started with "/".</li> </ul> </li> </ul>

Figure 5: Example details for dynamic non-conformance.  
Table 1: CATMA's performance statistics on multiple MSAs

Name	#LOC	# Services	# Detected Non-Conformances (static / dynamic)	Avg. Runtime (seconds)
Springboot-Microservice <sup>2</sup>	879	9	0 / 16	4.279
microservice <sup>3</sup>	3117	7	2 / 1	2.995
spring-petclinic-microservices <sup>4</sup>	3990	12	1 / 26	78.559
piggymetrics <sup>5</sup>	9977	17	3 / 11	53.932

for analysing the 4 selected MSAs (averaged over 10 executions per MSA). This evaluation allows us to understand if there is any benefit in using CATMA over a manual analysis. The results show that an analysis can be conducted within a matter of minutes. Zhou et al. [16] reported, that developers typically spend days in resolving issues within their MSA. Thus, CATMA can help developers reduce the time spent on debugging issues.

**User Study.** We conducted a small-scale user study to investigate CATMA's usefulness. We report an initial assessment of this user study based on a think-aloud interview setup with two participants. The participants were recruited from the lab of one of the authors and have no relation to the work done for CATMA. The participants got an introduction to MSAs and were allowed to interact with CATMA before the start of the interview. During the interview, we asked several questions that would provide us insights on what are the most useful elements presented in the output generated by CATMA. A complete transcript of the interview can be found at [4]. The following points summarize the most useful elements from CATMA's output: (1) the model-based visualization that shows where non-conformances are detected, (2) the set of possible interpretations providing the potential causes for the

<sup>2</sup><https://github.com/shabbirdwd53/Springboot-Microservice>

<sup>3</sup><https://github.com/ewolff/microservice>

<sup>4</sup><https://github.com/spring-petclinic/spring-petclinic-microservices>

<sup>5</sup><https://github.com/sqshq/piggymetrics>

**Table 2: Trade-off between the size and correctness.**

Avg. # Edges	Avg. # Nodes	Avg. Recall	Avg. Specificity	Avg. Balanced Accuracy
1	1	0.0	1.0	0.5
127	99	0.368	0.998	0.683
790	646	0.904	0.990	0.947
1982	1843	0.920	0.986	0.953
4714	4570	1.0	0.978	0.989

corresponding non-conformance, (3) the ability to jump from the runtime model (state machine) back to the source code, (4) static non-conformances provide insights on the security implication of the system, and (5) the type of the non-conformances: static non-conformances provide insights on the security implication.

**Correctness of Dynamic Models.** As the state machines approximate the provided input data, it is useful to understand the trade-off between the correctness and the size of the model as it could influence the detection of non-conformances; a small state machine generalizes too much and introduces inaccuracies, a large state machine captures all possible behavior but might be hard to understand and process. To evaluate this aspect, we use a technique similar to the one proposed by Walkinshaw et al. [13]. Table 2 presents the average results computed from a 10-fold cross-validation experiment. The size of a state machine is computed as the sum of states and edges. As expected, the results show that the correctness of a model drops when its size decreases. Furthermore, the results show that there is an optimal size for the models as the correctness does not significantly increase when a larger model is learned from data.

## 5 RELATED WORK

Several approaches in related literature combine static and dynamic analysis for architecture reconstruction of MSAs. *MicroArt* presented by Granchelli et al. [5], *MiSAR* presented by Alshuqayran et al. [1], and  *$\mu$ TOSCA* presented by Soldani et al. [12] all extract the list of microservices statically by parsing deployment files. Connections between them are detected dynamically by leveraging service discovery services that exist in the analyzed applications or by injecting different monitoring tools. *VMAWV* presented by Ma et al. [9] instead queries existing service discovery services to retrieve the list of services and uses static analysis to detect connections. While these approaches combine static and dynamic analysis, none of them compare complete architectural models obtained via the two techniques. Since our approach performs this comparison to identify non-conformances, we believe it to be novel in this regard.

## 6 CONCLUSION & FUTURE WORK

We present CATMA, a tool for automatically conducting conformance analysis of MSAs. It detects possible non-conformances by computing differences between a statically and a dynamically obtained architectural model of the MSA. Found non-conformances are visualized in an easily accessible way. Further, a set of possible interpretations is generated, showing the non-conformances' potential causes. In a preliminary evaluation, CATMA showed promising results in terms of performance as well as usability for the identification of non-conformances by developers. In our evaluation, CATMA identified a non-conformance in an open-source MSA on GitHub. A misconfiguration in the Hystrix<sup>6</sup> monitoring dashboard prevented

stream data from being visualized as it was intended in the implementation. With that, it is a good example of a non-conformance between intended and observed behavior of the MSA. We notified the developer and our fix was accepted<sup>7</sup>. Hence, CATMA has already shown its first –albeit small– impact on MSA.

As future work, the approach would benefit from more extensive validation activities concerning its usefulness and possible enhancements. Specifically, we plan an extended user study with developers of MSA in which they identify non-conformances with the help of CATMA. Further, we will investigate the feasibility of using other types of models as input and the detection capabilities of other non-conformances.

## ACKNOWLEDGMENTS

This work was partly funded by the European Union's Horizon 2020 program under grant agreement No. 952647 (AssureMOSS).

## REFERENCES

- [1] Nuha Alshuqayran, Nour Ali, and Roger Evans. 2018. Towards Micro Service Architecture Recovery: An Empirical Study. In *ICSA*.
- [2] Danilo Caivano, Pietro Cassieri, Simone Romano, and Giuseppe Scanniello. 2023. On the spread and evolution of dead methods in Java desktop applications: an exploratory study. *Empirical Software Engineering*.
- [3] Clinton Cao, Agathe Blaise, Sicco Verwer, and Filippo Rebecchi. 2022. Learning State Machines to Monitor and Detect Anomalies on a Kubernetes Cluster. In *ARES*.
- [4] Clinton Cao, Simon Schneider, Nicolás Díaz Ferreyra, Sicco Verwer, and Riccardo Scandariato. 2023. Appendix for 'CATMA: Conformance Analysis Tool for Microservice Applications'. <https://figshare.com/s/bce4718ac6f07b26742d>
- [5] Giona Granchelli, Mario Cardarelli, Paolo Di Francesco, Ivano Malavolta, Ludovico Iovino, and Amleto Di Salle. 2017. MicroART: A software architecture recovery tool for maintaining microservice-based systems. In *ICSAW 2017*.
- [6] Xiaofeng Guo, Xin Peng, Hanzhang Wang, Wanxue Li, Huai Jiang, Dan Ding, Tao Xie, and Liangfei Su. 2020. Graph-based trace analysis for microservice architecture understanding and problem diagnosis. In *ESEC/FSE*.
- [7] Abdelhakim Hannousse and Salima Yahiouche. 2021. Securing microservices and microservice architectures: A systematic mapping study. *Computer Science Review*.
- [8] Valentina Lenarduzzi and Annibale Panichella. 2021. Serverless Testing: Tool Vendors' and Experts' Points of View. *IEEE Software*.
- [9] Shang Pin Ma, I. Hsiu Liu, Chun Yu Chen, Jiun Ting Lin, and Nien Lin Hsueh. 2019. Version-Based Microservice Analysis, Monitoring, and Visualization. *APSEC*.
- [10] Simon Schneider, Tufan Özen, Michael Chen, and Riccardo Scandariato. 2023. microSecEnD: A Dataset of Security-Enriched Dataflow Diagrams for Microservice Applications. In *MSR*.
- [11] Simon Schneider and Riccardo Scandariato. 2023. Automatic extraction of security-rich dataflow diagrams for microservice applications written in Java. *Journal of Systems and Software*.
- [12] Jacopo Soldani, Giuseppe Muntoni, Davide Neri, and Antonio Brogi. 2021. The mTOSCA toolchain: Mining, analyzing, and refactoring microservice-based architectures. *Software: Practice and Experience*.
- [13] Neil Walkinshaw, Ramsay Taylor, and John Derrick. 2016. Inferring extended finite state machine models from software executions. *Empirical Software Engineering*.
- [14] Muhammad Waseem, Peng Liang, Mojtaba Shahin, Aakash Ahmad, and Ali Rezaei Nassab. 2021. On the Nature of Issues in Five Open Source Microservices Systems: An Empirical Study. In *EASE*.
- [15] Erik Whiting and Sharon Andrews. 2020. Drift and Erosion in Software Architecture: Summary and Prevention Strategies. In *ICISDM*.
- [16] Xiang Zhou, Xin Peng, Tao Xie, Jun Sun, Chao Ji, Wenhai Li, and Dan Ding. 2018. Fault Analysis and Debugging of Microservice Systems: Industrial Survey, Benchmark System, and Empirical Study. *IEEE Transactions on Software Engineering*.
- [17] Xiang Zhou, Xin Peng, Tao Xie, Jun Sun, Chao Ji, Dewei Liu, Qilin Xiang, and Chuan He. 2019. Latent error prediction and fault localization for microservice applications by learning from system trace logs. In *ESEC/FSE*.

<sup>6</sup><https://github.com/Netflix/Hystrix>

<sup>7</sup><https://github.com/ewolff/microservice/pull/30>